


```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

```
In [7]: X_driving_hours = [10, 9, 2, 15, 10, 16, 11, 16]
y_risk_score = [95, 80, 10, 50, 45, 98, 38, 93]

plt.plot(X_driving_hours, y_risk_score, "o")
plt.xlabel("Number of Hours spent driving (x)")
plt.ylabel("Risk score on a scale of 0-100 (y)")

plt.show()
```



```
In [8]: def mean(values):
return sum(values)/float(len(values))
```

```
In [9]: def variance(values, mean):
return sum([(x-mean)**2 for x in values])
```

```
In [10]: def covariance(x, x_mean, y, y_mean):
cv = 0.0

for i in range(len(x)):
cv += (x[i] - x_mean) * (y[i] - y_mean)

return cv
```

```
In [16]: def find_coefficients(x, y):
x_mean = mean(x)
y_mean = mean(y)

b1 = covariance(x, x_mean, y, y_mean) / variance(x, x_mean)
b0 = y_mean - b1 * x_mean

return [b0, b1]
```

```
In [23]: b0, b1 = find_coefficients(X_driving_hours, y_risk_score)
print("Coefficients are :\nb0 = %.3f \nb1 = %.3f" % (b0, b1))

Coefficients are :
b0 = 12.585
b1 = 4.588
```

```
In [27]: print("The equation of the best fit line (y = b0 + b1x) is : y = %.3f + %.3fx" % (b0, b1))

The equation of the best fit line (y = b0 + b1x) is : y = 12.585 + 4.588x
```

```
In [29]: # Checking if coefficients are correct or not

from sklearn.linear_model import LinearRegression

X = np.array(X_driving_hours).reshape(-1, 1)
y = np.array(y_risk_score).reshape(-1, 1)

lr = LinearRegression().fit(X, y)
```


```
In [30]: lr.coef_
Out[30]: array([[4.58789861]])
```

```
In [31]: lr.intercept_
Out[31]: array([12.58462796])
```

```
In [43]: def abline(intercept, slope):
axes = plt.gca()
xvals = np.array(axes.get_xlim())
yvals = intercept + slope * xvals

plt.plot(xvals, yvals, "--")
```

```
In [44]: plt.plot(X_driving_hours, y_risk_score, "o")
abline(b0, b1)
```



In [50]: # Prediction

```
def linear_regression(X_train, y_train, X_test):
    predictions = list()

    b0, b1 = find_coefficients(X_train, y_train)

    for row in X_test:
        y_pred = b0 + b1 * row
        predictions.append(y_pred)

    return predictions
```

In [52]: n = int(input("Enter the number of predictions required: "))

```
res = []

for i in range(n):
    ele = int(input())
    res.append(ele)

predictions = linear_regression(X_driving_hours, y_risk_score, res)

predictions

Enter the number of predictions required: 3
3
7
14
```

Out[52]: [26.348323793949305, 44.699918233851186, 76.81520850367949]

In [63]: # Calculate RMSE

```
from math import sqrt

def calc_rmse(actual, prediction):
    sum_err = 0.0

    for i in range(len(actual)):
        pred_err = prediction[i] - actual[i]
        sum_err += (pred_err ** 2)

    mean_err = sum_err / float(len(actual))

    return sqrt(mean_err)
```

In [77]: # Splitting the data

```
from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X_driving_hours, y_risk_score, test_size=1/3, random_state=0)

pred = linear_regression(X_train, y_train, X_test)
rmse = calc_rmse(y_test, pred)


print("Root Mean Squared Error (RMSE): ", rmse)

Root Mean Squared Error (RMSE): 29.547856861944695
```

In [78]: plt.plot(X\_train, y\_train, "bo")
plt.plot(X\_test, y\_test, "ro")
plt.plot(X\_test, pred, "y")

```
plt.plot()

abline(b0, b1)
plt.show()
```



In [ ]: